

# workshop04tasks

March 21, 2021

## 1 Welcome to Week 4: 1. Clustering

- 

1.1 Task to be discussed in this Workshop are:

- 

1.2 Demo for K-means clustering

- 

1.3 Task 1.1 Perform K-means on a real dataset

- 

1.4 Task 1.2 (Optional) Try PCA for dimensionality reduction.

- 

1.5 Task 1.3: Perform agglomerative clustering on this data set

```
[1]: import math
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt

#from sklearn.datasets.samples_generator import make_blobs
from sklearn.datasets import make_blobs


from sklearn.cluster import KMeans
from sklearn.metrics import davies_bouldin_score, adjusted_rand_score

from sklearn.metrics import pairwise_distances
from scipy.cluster.hierarchy import linkage, dendrogram, cut_tree
from scipy.spatial.distance import pdist
```

```
[ ]:
```

```
[2]: # Generate an example 2-dimensional datasets containing 4 clusters for a demo
X, y = make_blobs(n_samples=300, centers=4, cluster_std=0.60, random_state=0)

# Visualize the data
plt.scatter(X[:,0], X[:,1])
plt.show()
```



```
[3]: # Create a K-means clustering model with k=4, and k-means++ as the
      ↪ initialization strategy
kmeans = KMeans(n_clusters=4, init='k-means++', max_iter=300, random_state=0)

# Perform clustering by fitting the model with the data
kmeans.fit(X)

# We can explore the parameters learned from the data.
# What's the cluster center?
print('\n Cluter center: \n', kmeans.cluster_centers_)

# What's is overall distorion (inertia)?
print('\n Overall distortion: \n', kmeans.inertia_)
# Average standard distortion
print('\n Average distortion: \n', math.sqrt(kmeans.inertia_/X.shape[0]))
```

```
Cluter center:
[[ 1.98258281  0.86771314]
```

```
[ 0.94973532  4.41906906]
[-1.37324398  7.75368871]
[-1.58438467  2.83081263]]
```


Overall distortion:  
212.00599621083472

Average distortion:  
0.8406465690384489

```
[4]: # Predict the cluster for each data instance. This step can be combined with
      ↪ the last one by using kmeans.fit_predict(X)
y_pred = kmeans.predict(X)

# Visualize the cluster centers to explore how the clustering result looks like
colors = ['blue', 'yellow', 'green', 'cyan']
for i, color in enumerate(colors):
    plt.scatter(X[y_pred == i, 0], X[y_pred == i, 1], c=color)

plt.scatter(kmeans.cluster_centers_[ :, 0], kmeans.cluster_centers_[ :, 1],
            ↪ marker='x', lw=2, c='red', s=100)
plt.show()
```



```
[5]: # How about other k values? Which k value should we choose to get the optimal
      ↪ clustering? Let's vary k from 1 to 10 and see how the distortion changes.
distortions = []
```


```

for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300,
        random_state=0)
    kmeans.fit(X)
    distortions.append(kmeans.inertia_)

# Plot the relationship between the distortion and k. Then, we can have the
    Elbow method to help identify a good value for k.
plt.plot(range(1, 11), distortions)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('Overall distortion')
plt.scatter(4, distortions[4], marker='x', lw=2, c='red', s=100)
plt.show()

# It can be seen that k=4 is reasonably good

```



```

[6]: # Let's evaluate the learned model with other quality criteria

# Internal evaluation, davies bouldin score (the lower, the better)
db_scores = []
db_scores_std = []

```

```

# External evaluation, adjusted rand index (the higher, the better)
ar_scores = []
ar_scores_std = []

# Inertia (average standardized)
inertia = []
inertia_std = []

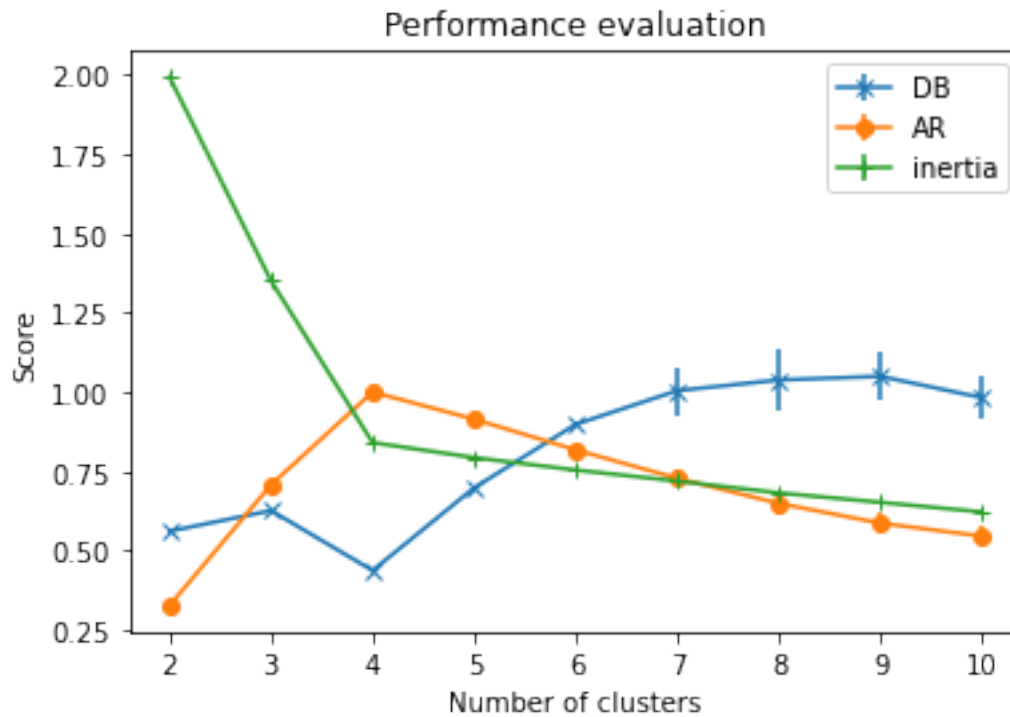
for i in range(2, 11):
    # Multiple runs for stable indicators
    db_scores_tmp = []
    ar_scores_tmp = []
    inertias_tmp = []

    n_iteration=5
    for j in range(0, n_iteration):
        # Note: 'random_state' parameter should be set as default or None
        #kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300,
↪n_init=10)
        kmeans = KMeans(n_clusters=i, init='random', max_iter=300)
        kmeans.fit(X)
        labels = kmeans.labels_
        db_scores_tmp.append(davies_bouldin_score(X, labels))
        ar_scores_tmp.append(adjusted_rand_score(labels, y_pred))
        inertias_tmp.append(math.sqrt(kmeans.inertia_/X.shape[0]))

    db_scores.append(np.mean(db_scores_tmp))
    db_scores_std.append(np.std(db_scores_tmp))
    ar_scores.append(np.mean(ar_scores_tmp))
    ar_scores_std.append(np.std(ar_scores_tmp))
    inertia.append(np.mean(inertias_tmp))
    inertia_std.append(np.std(inertias_tmp))

# Plot the relationship between the davies bouldin score and k
plt.errorbar(range(2, 11), db_scores, yerr=db_scores_std, marker='x',
↪label='DB')
plt.errorbar(range(2, 11), ar_scores, yerr=ar_scores_std, marker='o',
↪label='AR')
plt.errorbar(range(2, 11), inertia, yerr=inertia_std, marker='+',
↪label='inertia')
plt.title('Performance evaluation')
plt.xlabel('Number of clusters')
plt.ylabel('Score')
plt.legend(loc='best')
plt.show()

```




[7]: *# Try agglomerative clustering on this dataset, and visualise the hierarchy.*

```
dist = pdist(X, 'euclidean')
linkage_matrix = linkage(dist, method = 'complete')
```

```
plt.figure(figsize=(15,7))
dendrogram(linkage_matrix)
plt.show()
```

*# It can be seen that the clustering structure contains four main clusters, □  
↪ complying with the data.*



## 1.6 Task Clustering on the MNIST data set

```
[8]: # Now, let's work on a real dataset. See detailed information for the dataset:
      ↪ https://en.wikipedia.org/wiki/MNIST\_database.

      #Load the data. Orignal data set has been processed (downsampled) to facilitate
      ↪ your data analysis
      raw_data = pd.read_csv("/home/dalia/Master of Information Technology in
      ↪ Cybersecurity/Session 3/COMP8325 Artificial Intelligence/week_4/Week 4
      ↪ Workshop -20210321/data/mnist-0.1.csv")
      print('\n data size: (%d, %d)\n' % raw_data.shape)

      # Specifying features and target attribute
      X = raw_data.drop(['Label'], axis='columns')

      # Pre-processing with standardization
      from sklearn import preprocessing
      scaler = preprocessing.MinMaxScaler()
      X_data = X.values
      X_scaled = scaler.fit_transform(X_data)
      X = X_scaled

      y = raw_data['Label'].values
```

data size: (5243, 785)

```
[9]: # K-means clustering model
model = KMeans(n_clusters=2)
model.fit(X)
print('\n cluster means: \n', model.cluster_centers_)
print('\n inertia: %f'% model.inertia_)
print('\n average inertia: %f\n' % math.sqrt(model.inertia_/y.size))
```

```
cluster means:
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
inertia: 214710.126187
```

```
average inertia: 6.399357
```

```
[10]: # Evaluation (internal)
labels = model.labels_
scores=davies_bouldin_score(X, labels)
print('\n davies_bouldin_score: %f\n' % scores)
```

```
davies_bouldin_score: 2.722454
```

```
[11]: # Task 1.1 Try k from 2 to 10 to determine which is the best value w.r.t.
      ↪davies_bouldin_score, plot the relationship between the davies_bouldin_score
      ↪and k


# Compute davies bouldin score

distortions = []
for i in range(2, 11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(X)
    labels = kmeans.labels_
    distortions.append(davies_bouldin_score(X, labels))

# Plot the relationship between the distortion and k

plt.plot(range(2, 11), distortions)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('Overall distortion')
plt.scatter(4, distortions[4], marker='x', lw=2, c='red', s=100)
plt.show()
```





```
[12]: # Task 1.2 (Optional) Given that this is a high dimensional data. It might be
      ↪ good to reduce the dimension first.
      # PCA can be used for this purpose. Try some reduced dimensionality, e.g., math.
      ↪ sqrt(X.shape[1]). Try this for different k values with plotting.

      # dimension reduction
      from sklearn.decomposition import PCA
      pca = PCA(n_components=int(math.sqrt(X.shape[1])))

      # Compute davies bouldin score


      X_reduced = pca.fit(X).transform(X)
      X = X_reduced
      print('\n After dimension reduction: (%d, %d)\n' % X.shape)

      distortions = []
      for i in range(2, 11):
          kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
              ↪ random_state=0)
          kmeans.fit(X)
          labels = kmeans.labels_
          distortions.append(davies_bouldin_score(X, labels))
```

```
# Plot the relationship between the score and k

plt.plot(range(2, 11), distortions)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('Overall distortion')
plt.scatter(4, distortions[4], marker='x', lw=2, c='red', s=100)
plt.show()
```


After dimension reduction: (5243, 28)



[15]: *# Task 1.3 Try to perform agglomerative clustering on the dataset, and  
↪ visualise the hierarchy.*

```
dist = pdist(X, 'euclidean')
linkage_matrix = linkage(dist, method = 'complete')

plt.figure(figsize=(15,7))
dendrogram(linkage_matrix)
plt.show()
```



[ ]: